

# Deep Analysis: Cartedo Simulation Adaptation Framework

Date: 2026-01-16 Total Lines of Code: 23,300 Total Files: 42

---

## Part 1: Alignment & Adaptation Logic (How It Actually Works)

### The Complete Pipeline Flow

#### STEP 1: INPUT

|  
v

#### STEP 2: FACTSHEET EXTRACTION (Gemini)

- Extract company name, manager, industry
- Generate poison list (terms to avoid)
- Generate KLOs for TARGET scenario
- Generate industry\_context (KPIs, terminology)

|  
v

#### STEP 3: SHARDER

- Split JSON into 15 shards
- Lock 2 shards (workspace\_ids, scenario\_options)
- Hash each shard for change detection

|  
v

#### STEP 4: PARALLEL SHARD ADAPTATION (Gemini)

- Each unlocked shard sent to LLM
- Factsheet injected into every prompt
- RAG context added per shard type

|  
v

#### STEP 5: MERGE SHARDS

- Combine 15 shards back into single JSON
- Run cleanup\_merged\_json()

|  
v

#### STEP 6: ALIGNMENT CHECKER (GPT 5.2)

- 9 parallel LLM checks
- Calculate weighted score
- If score < 95% -> trigger fixer

|

v

#### STEP 7: ALIGNMENT FIXER (GPT 5.2)

- Phase 1: Fast fixes (string replacement)
- Phase 2: Batched LLM fix (questions, content, resources)

|  
v

#### STEP 8: VALIDATION AGENT (GPT 5.2)

- 8 Critical checks (blocking)
- 6 Flagged checks (non-blocking)

|  
v

#### STEP 9: FINISHER

- Compute weighted compliance score
- Decision: PASS / FAIL\_RETRY / FAIL\_HUMAN

|  
v

#### STEP 10: OUTPUT

---

## Step-by-Step Breakdown

STEP 1-2: Factsheet Extraction (The Foundation) File: `src/utils/gemini_client.py` (846 LOC)

What happens:

```
# Line 234-237 in adaptation_engine.py
global_factsheet = await extract_global_factsheet(
    source_scenario=source_scenario, # "HR Hiring at Summit Innovations"
    target_scenario=target_scenario, # "Market Entry for EcoChic Threads"
)
```

Output structure:

```
{
  "company": {
    "name": "EcoChic Threads",
    "industry": "Sustainable Fashion Retail",
    "type": "DTC Startup"
  },
  "reporting_manager": {
    "name": "Maya Sharma",
    "role": "Director of Market Strategy",
    "email": "maya.sharma@ecochic.com",
    "gender": "female"
  },
  "poison_list": [
    "Elizabeth Carter", "Summit Innovations", "HR",
    "hiring", "candidate", "interview", "recruitment"
  ],
  "klos": [
    {"id": "klo1", "outcome": "Analyze market opportunity using TAM/SAM/SOM"},
    {"id": "klo2", "outcome": "Evaluate internal capabilities and resource fit"},
    {"id": "klo3", "outcome": "Develop go/no-go recommendation with risk assessment"}
  ],
  "industry_context": {
    "kpis": ["CAC", "LTV", "Gross Margin", "Market Share"],
    "terminology": ["market entry", "competitive analysis", "PESTEL"],
    "wrong_terms": ["HR metrics", "hiring KPIs", "recruitment funnel"]
  }
}
```

**PROBLEM IDENTIFIED:** The poison\_list is generated from the TARGET scenario prompt, but the prompt may not contain explicit domain terms like "HR", "hiring", "candidate". These are domain-IMPLIED terms that the LLM doesn't automatically include.

---

STEP 3: Sharder (Divide & Conquer) File: `src/stages/sharder.py` (407 LOC)

What happens:

```
# Splits JSON into 15 shards:
SHARD_DEFINITIONS = [
  "workspace_ids",           # LOCKED - never modified
  "scenario_options",        # LOCKED - never modified
  "lesson_information",      # Contains manager emails
  "overview",                # Scenario overview text
  "guidelines",              # Student guidelines
  "simulation_flow",          # Core activities + resources
  "submission_questions",    # Assessment questions
  "rubric",                  # Grading criteria
  "assessment_criterion",   # KLO definitions
```

```

"emails",           # Manager communications
"characters",      # Personas
"resources",       # Learning materials
"workplace_scenario", # Scenario narrative
"activities",      # Task definitions
"stage_definitions", # Workflow stages
]

```

**PROBLEM IDENTIFIED:** KLOs are in `assessment_criterion` shard but questions are in `simulation_flow` shard. When processed independently, there's no cross-shard awareness of KLO-question alignment.

---

**STEP 4-5: Parallel Adaptation & Merge** File: `src/stages/adaptation_engine.py` (846 LOC)

What happens:

```

# Each shard gets this prompt structure:
prompt = f"""
## GLOBAL CONTEXT (MANDATORY):
{factsheet_json}

## POISON LIST (NEVER USE THESE):
{poison_list}

## RAG CONTEXT (Industry Examples):
{rag_context}

## TASK:
Adapt this {shard_name} content from {source} to {target}.

## INPUT SHARD:
{shard_json}
"""

```

**PROBLEM IDENTIFIED:** Each shard is adapted INDEPENDENTLY. Even though factsheet is injected, the LLM doesn't see the actual KLO definitions when adapting questions in `simulation_flow` shard.

---

**STEP 6: Alignment Checker (Quality Gate)** File: `src/stages/alignment_checker.py` (1,472 LOC)

What happens:

```

# 9 parallel checks run using asyncio:
check_coroutines = [
    self._check_reporting_manager_consistency(), # R1
    self._check_company_consistency(),          # R2
    self._check_poison_terms(),                 # R3
    self._check_klo_to_questions(),             # R4 <- FAILING (90%)
    self._check_klo_to_resources(),              # R5 <- FAILING (88%)
    self._check_scenario_to_resources(),        # R6 <- FAILING (90%)
    self._check_role_to_tasks(),                # R7 <- FAILING (90%)
    self._check_klo_alignment(),                # R8 <- FAILING (90%)
    self._check_scenario_coherence()           # R9 <- FAILING (90%)
]

```

Score Calculation:

```

# Line 322-324
if results:
    overall_score = sum(r.score for r in results) / len(results)
# Example: (96+97+100+90+88+90+90+90+90) / 9 = 92.33%

```

**PROBLEM IDENTIFIED:** The LLM checks run against the MERGED JSON, so they CAN see cross-shard relationships. But by this point, the content is already wrong - the checker can only DETECT issues, not prevent them.

---

**STEP 7: Alignment Fixer (The Repair Agent)** File: src/stages/alignment\_fixer.py (1,782 LOC)

What happens:

```
# Phase 1: Fast fixes (no LLM)
if rule_id == "reporting_manager_consistency":
    # String replacement: find/replace manager names

# Phase 2: Batched LLM fix
if llm_fix_rules:
    batched_result = await self._fix_all_batched(
        json_data=fixed_json,
        failed_rules=llm_fix_rules,
        factsheet=global_factsheet,
    )
```

The Bug (Lines 391-399):

```
# Looks for questions HERE:
submission_questions = topic_data.get("submissionQuestions", [])

# But questions are actually in simulationFlow[].data.submissionQuestions!
# This is why fixes=0 - it can't find the questions to fix
```

**PROBLEM IDENTIFIED:** Alignment fixer searches for `submissionQuestions` at the TOP LEVEL, but they're nested inside `simulationFlow[] .data.submissionQuestions`. The fixer returns `fixes=0` because it finds an empty list.

---

Why Alignment Score is Stuck at 92.33%

Rule	Score	Why It Fails
R4: KLO-Questions	90%	Questions are duplicated across KLOs, not KLO-specific
R5: KLO-Resources	88%	Resources truncated, missing financial model
R6: Scenario-Resources	90%	Resources still contain HR terminology
R7: Role-Tasks	90%	Tasks don't map to Junior Consultant role
R8: KLO-Task	90%	Activities not linked to specific KLOs
R9: Coherence	90%	"optimization" vs "go/no-go" mismatch

Weighted Average:  $(96+97+100+90+88+90+90+90+90) / 9 = 92.33\%$

---

## Part 2: Agent Clashes & Conflicts

### CLASH 1: Alignment Fixer vs Semantic Fixer

Aspect	Alignment Fixer	Semantic Fixer
File	alignment_fixer.py (1,782 LOC)	fixers.py (1,612 LOC)
Triggered By	Alignment score < 95%	Validation failures
Fixes	KLO mapping, question alignment	Entity replacement, KPIs
Model	GPT 5.2	GPT 5.2
Input	alignment_report	validation_report

**THE CLASH:** Both agents try to fix “HR terminology” but neither succeeds because:  
 - Alignment Fixer: Can’t find questions (wrong path)  
 - Semantic Fixer: Runs AFTER alignment fixer, by then alignment score is locked

**FIX NEEDED:** Either merge the two fixers OR ensure alignment fixer runs BEFORE semantic fixer with correct paths.

---

### CLASH 2: Adaptation Engine vs Alignment Checker

Aspect	Adaptation Engine	Alignment Checker
File	<code>adaptation_engine.py</code> (846 LOC)	<code>alignment_checker.py</code> (1,472 LOC)
Model	Gemini	GPT 5.2
When	BEFORE shards merge	AFTER shards merge
Context	Per-shard only	Full JSON

**THE CLASH:** - Adaptation uses Gemini with per-shard context - Alignment uses GPT 5.2 with full JSON context - When Gemini makes a mistake, GPT detects it but can’t fix it (different model, different context)

**FIX NEEDED:** Either use same model for both OR pass full KLO context to Gemini during adaptation.

---

### CLASH 3: Check Definitions vs Alignment Rules

Source	Threshold	What It Checks
<code>check_definitions.py</code> C7	95%	KLO Preservation
<code>alignment_checker.py</code> R4	95%	KLO-Questions
<code>alignment_checker.py</code> R5	95%	KLO-Resources

**THE CLASH:** Three different places check “KLO alignment” with slightly different logic: - C7 checks if KLOs exist and are preserved - R4 checks if questions map to KLOs - R5 checks if resources support KLOs

A run can pass C7 but fail R4/R5, creating confusion about what “KLO alignment” means.

**FIX NEEDED:** Consolidate KLO checks into single source of truth.

---

### CLASH 4: Finisher Thresholds vs Alignment Thresholds

Stage	Threshold	Pass Condition
Alignment Checker	95%	<code>overall_score &gt;= 0.95</code>
Finisher	80% blocker, 70% overall	<code>blocker_pass_rate &gt;= 0.8 AND overall &gt;= 0.7</code>
Dashboard	95%	<code>acceptance_threshold = 0.95</code>

**THE CLASH:** - Alignment says 92.33% = FAIL - Finisher says 100% compliance = PASS - Dashboard says 95% = FAIL

A run can PASS finisher but FAIL alignment, creating confusion about release readiness.

**FIX NEEDED:** Align all thresholds to single standard (95%).

---

## Part 3: Detailed Agent Analysis

### STAGE 1: Adaptation Engine

File: src/stages/adaptation\_engine.py Lines of Code: 846 Model: Gemini 2.0 Flash

**One-Line Summary:** Generates initial adapted JSON by extracting factsheet, splitting into shards, and running parallel LLM adaptation with RAG context.

#### Key Functions:

Function	Lines	What It Does
extract_klos_from_json()	50-91	Extracts KLOs from multiple possible locations in source JSON
format_klos_for_prompt()	93-124	Formats KLOs for injection into adaptation prompts
AdaptationEngine.adapt()	152-400	Main entry point - orchestrates full adaptation
_adapt_shard()	401-500	Adapts single shard with factsheet + RAG context
_get_shard_rag_context()	501-550	Retrieves shard-specific examples from ChromaDB

#### Data Flow:

```
Input JSON
|
v
extract_global_factsheet() --> factsheet
|
v
Sharder.shard() --> 15 shards
|
v
[Parallel] _adapt_shard() x 13 --> adapted shards
|
v
merge_shards() --> adapted JSON
```

**Critical Dependencies:** - `gemini_client.py` for LLM calls - `sharder.py` for JSON splitting - `prompts.py` for prompt templates

---

### STAGE 2: Sharder

File: src/stages/sharder.py Lines of Code: 407 Model: None (pure Python)

**One-Line Summary:** Splits simulation JSON into 15 semantically independent shards with lock states and hash tracking for change detection.

#### Key Functions:

Function	Lines	What It Does
Sharder.shard()	50-150	Splits JSON into shards based on SHARD_DEFINITIONS
Sharder.merge_shards()	151-250	Combines shards back into single JSON
cleanup_merged_json()	251-350	Post-merge cleanup (URL fix, placeholder removal)
compute_shard_hash()	351-400	SHA256 hash for change detection

#### Shard Lock States:

```
class LockState(Enum):
    UNLOCKED = "unlocked"      # Can be adapted
    LOCKED = "locked"          # Never sent to LLM
    LOCKED_AFTER_FIX = "locked_after_fix" # Locked after structural fix
```

---

### STAGE 3: Alignment Checker

File: `src/stages/alignment_checker.py` Lines of Code: 1,472 Model: GPT 5.2

**One-Line Summary:** Validates cross-shard semantic alignment using 9 parallel LLM checks with weighted scoring against 95% threshold.

#### Key Functions:

Function	Lines	What It Does
<code>AlignmentChecker.check()</code>	228-351	Main entry - runs all 9 checks in parallel
<code>_check_reporting_manager_consistency()</code>	108-185	Verifies manager name/email/role consistency
<code>_check_company_consistency()</code>	487-560	Verifies company name consistency
<code>_check_poison_terms()</code>	562-650	Scans for source scenario leakage
<code>_check_klo_to_questions()</code>	652-750	Maps questions to KLOs
<code>_check_klo_to_resources()</code>	752-850	Maps resources to KLOs
<code>_check_scenario_to_resources()</code>	852-950	Verifies scenario-resource alignment
<code>_check_role_to_tasks()</code>	952-1050	Verifies role-task alignment
<code>_check_klo_alignment()</code>	1052-1150	Verifies overall KLO structure
<code>_check_scenario_coherence()</code>	1152-1250	Verifies narrative coherence

#### Scoring Guidelines (Lines 358-379):

0.95-1.0: Perfect  
0.85-0.94: Excellent  
0.75-0.84: Good  
0.65-0.74: Acceptable  
0.50-0.64: Needs Work  
Below 0.50: Poor

---

### STAGE 3B: Alignment Fixer

File: `src/stages/alignment_fixer.py` Lines of Code: 1,782 Model: GPT 5.2

**One-Line Summary:** Fixes alignment-specific issues using fast string replacement for consistency and batched LLM calls for KLO/resource alignment.

#### Key Functions:

Function	Lines	What It Does
<code>AlignmentFixer.fix()</code>	147-250	Main entry - orchestrates fix phases
<code>_fix_manager_consistency()</code>	500-600	Fast string replacement for manager names
<code>_fix_company_consistency()</code>	601-700	Fast string replacement for company names
<code>_fix_all_batched()</code>	252-362	Batched LLM fix for all KLO/resource issues
<code>_fix_questions_batch()</code>	364-500	Fixes submission questions for KLO alignment
<code>_fix_content_batch()</code>	700-900	Fixes guidelines/overview/scenario text
<code>_fix_resources_batch()</code>	901-1100	Fixes resources for industry alignment

#### THE BUG (Line 396):

```
# WRONG: Looks at top level
submission_questions = topic_data.get("submissionQuestions", [])

# CORRECT: Should look in simulationFlow
for stage in topic_data.get("simulationFlow", []):
    stage_questions = stage.get("data", {}).get("submissionQuestions", [])
```

---

## STAGE 4: Fixers (Structural + Semantic)

File: src/stages/fixers.py Lines of Code: 1,612 Model: GPT 5.2

**One-Line Summary:** Applies surgical JSON patches for structural (shape) and semantic (meaning) fixes using JSON Pointer operations.

### Key Functions:

Function	Lines	What It Does
StructuralFixer.fix()	100-300	Fixes schema violations (missing keys, wrong types)
SemanticFixer.fix()	301-600	Fixes content issues (entities, KPIs, terminology)
_identify_fixes()	601-800	Uses LLM to identify exact JSON paths to fix
_apply_patches()	801-1000	Applies JSON Pointer patches to JSON

### Fix Types:

```
class FixType(Enum):
    STRUCTURAL = "structural" # Shape only
    SEMANTIC = "semantic"     # Meaning only
```

## STAGE 5: Validation Agent

File: src/validation/validation\_agent.py Lines of Code: 360 Model: GPT 5.2

**One-Line Summary:** Produces human-readable validation reports for PM/QA decision-making with 8 critical (blocking) and 6 flagged (non-blocking) checks.

### Key Functions:

Function	Lines	What It Does
ValidationAgent.validate()	50-150	Main entry - runs all checks
_run_critical_checks()	151-220	Runs 8 blocking checks
_run_flagged_checks()	221-280	Runs 6 non-blocking checks
_generate_report()	281-350	Generates structured report

## STAGE 6: Finisher

File: src/stages/finisher.py Lines of Code: 392 Model: None (pure Python)

**One-Line Summary:** Computes weighted compliance score from validation results and makes final pass/fail decision.

### Key Functions:

Function	Lines	What It Does
Finisher.run()	74-150	Main entry - computes compliance
_compute_compliance_score()	151-250	Calculates weighted score
_decide_status()	251-300	Returns PASS/FAIL_RETRY/FAIL_HUMAN

### Pass Conditions:

```
# Line 41-43
@property
def passed(self) -> bool:
    return self.blocker_pass_rate >= 0.8 and self.overall_score >= 0.70
```

## Part 4: 10 Deep Questions for Seniors

### Question 1: Alignment Fixer Path Bug

“The Alignment Fixer searches for questions at `topic_data.submissionQuestions`, but they’re actually in `simulationFlow[] .data.submissionQuestions`. This bug has been present for 4+ runs. Should we add schema documentation that maps where each content type lives to prevent similar issues?”

### Question 2: Cross-Model Repair Strategy

“Gemini adapts the content, GPT detects issues, then GPT tries to fix Gemini’s mistakes. Have we measured if this cross-model repair actually works, or does GPT just introduce different errors?”

### Question 3: Threshold Inconsistency

“We have three different pass thresholds: Alignment (95%), Finisher (70-80%), Dashboard (95%). Which is the actual release gate? A run can PASS finisher but FAIL alignment.”

### Question 4: Poison List Generation

“The poison list misses domain-implied terms like ‘HR’, ‘hiring’, ‘candidate’ (185 occurrences). Should we maintain a static domain-to-terms mapping, or ask the LLM to infer domain vocabulary?”

### Question 5: Shard Isolation vs Cross-Shard Alignment

“KLOs and questions are in different shards, processed independently. How do we maintain KLO-question alignment when each shard is adapted without seeing the other?”

### Question 6: Duplicate Check Systems

“Both Alignment Checker and Validation Agent check ‘KLO alignment’ with different logic. Should we consolidate into single source of truth?”

### Question 7: Resource Truncation Root Cause

“Resources are truncated mid-content. Is this (a) LLM output limit, (b) JSON parsing error, (c) context window exceeded, or (d) prompt doesn’t specify minimum length?”

### Question 8: Human Escalation Threshold

“After 4 consecutive failures (FAILED14-20) with same issues, no human escalation triggered. What should be the threshold for automatic human review?”

### Question 9: Retry Strategy

“`MAX_ATTEMPTS=1` for speed, but we’re stuck at 92.33%. Should we increase retries specifically for alignment-critical fixes that are close to threshold?”

### Question 10: Validation-Fixer Loop

“Validation agent produces detailed failure report, but fixers don’t consume it directly - they re-run their own checks. Why aren’t we passing validation failures to the specific fixer agent?”

---

## Part 5: One-Line Summary

“The 23,300 LOC pipeline achieves 92.33% alignment but has 4 critical clashes: (1) alignment fixer can’t find questions due to wrong path, (2) cross-model repair (Gemini->GPT) may not be effective, (3) threshold inconsistency between stages (95% vs 70%), and (4) shard isolation breaks cross-shard KLO alignment.”